

WGU D208 TASK 2 REV 3 - MATTINSON

Logistic Regression Using Churn Data

Mike Mattinson

Master of Science, Data Analytics, WGU.edu

D208: Predictive Modeling

Dr. Keiona Middleton

October 8, 2021

Abstract. This paper provides the results of a logistic regression analysis conducted on a customer dataset in partial fulfillment of WGU's D208 Predictive Analysis class requirements. The dataset represents 10,000 rows of customer data for a typical services company. There are fifty (50) attributes for each customer. The provided dataset was mostly clean and ready to use, however, some few additional data cleaning steps were completed prior to running the predictive analysis. The predictive analysis includes both an initial model using all the predictor variables and a final model using a reduced set of predictor variables. The final model includes both numerical and categorical predictor variables. P-values and multi-collinearity were used to select the features used in the final model. The principal research question "how to predict customer churn with high confidence using as few predictor variables as possible" was determined (93% correct predictions) using fourteen (14) of the original attributes. The analysis was conducted in a Python environment using a Jupyter notebook. The Jupyter notebook includes both code and discussion of the analysis. Key words: Churn. Regression. Logistic Regression. Primary data set: clean_churn.csv, the initial set has 10,000 records with 50 attributes.

Partial Reuse. A large portion of this notebook is re-used from Task 1 with the notable exception that this is a Logistic Regression instead of the Multiple Regression of Task 1. The same data was used in both tasks. Similar formatting was used in both tasks, as I am trying to be consistent with my submissions and preparation for future thesis work.

In addition, I found a pretty good example of logistic regression on "towards data science" website by Susan Li (2021), a lot of the flow and code was referenced and incorporated into this notebook.

Mattinson, M. (2021, September). WGU D208 TASK 1 REV 8 - MATTINSON.

Retrieved from: wgu.edu

Li, S. (2021, September). Building a Logistic Regression in Python, Step by Step.

Retrieved from: <https://towardsdatascience.com/building-a-logistic-regression-in-python-step-by-step-becd4d56c9c8>

Apply Custom Notebook Styles. Apply custom .css styles to the notebook. The 'custom.css' file is located in the project folder.

```
In [1]: # Styling notebook with custom css
from IPython.core.display import HTML
HTML(open('custom.css', "r").read())
```

Out[1]:

```
<div class="part">Research Question</div>
```

A1. RESEARCH QUESTION

Primary Research Question. A typical services company's revenue is maximized based on the total number of customers and how much each of those customers pay for those services. If the company charges too much, then the customer may stop the service, this is known as churn. If the company charges too little, then it will not maximize its revenue. This analysis will attempt to predict the probability of a customer's churn (dependent variable is 'Churn' which is a binary categorical data) using

logistic regression with high degree of accuracy based on a minimum set of predictor variables. The final set of predictor variables should include both numeric (e.g., Tenure, Child, and Income, etc.) and categorical data (e.g., Techie, Gender, and Internet Service type, etc.).

A2. OBJECTIVES AND GOALS

Data Preparation. Data Preparation objectives are addressed in Part III below and include the following:

- A. Convert categorical data.
- B. Mitigate missing data.
- C. Select data required for the analysis.
- D. Remove data deemed unnecessary.
- E. Explore data.
- F. Visualize data.
- G. Provide copy of final data.

Model Analysis. Model Analysis objectives are addressed in Part IV below and include the following:

- A. Eliminate predictor variables with high p-values.
- B. Eliminate predictor variables with high degree of multicollinearity.
- C. Create initial model using all the data.
- D. Refine model using a reduced set of the data.
- E. Summarize results.
- F. Ensure independent and dependent variables are linear.
- G. Ensure independent variables are not highly collinear
- H. Ensure final model residuals are normally distributed.

PART I: METHOD JUSTIFICATION

B1. ASSUMPTIONS

Assumptions. According to Massaron and Boschetti (2019), the logistic regression analysis is based on the following

assumptions:

- A. **Binary Dependent Variable.** Binary logistic regression requires the dependent variable to be binary.
- B. **Desired Outcome.** For a binary regression, the factor level 1 of the dependent variable should represent the desired outcome.
- C. **Only Meaningful Variables.** Only the meaningful variables should be included.
- D. **Multi-Collinearity.** The independent variables should be independent of each other. That is, the model should have little or no multicollinearity.
- E. **Independent Variable Linear to Log Odds.** The independent variables are linearly related to the log odds.
- F. **Large Sample Size.** Logistic regression requires large sample sizes.

Massaron, L., Boschetti, A. (2019). Regression Analysis with Python. Retrieved from:
<https://www.packtpub.com/product/regression-analysis-with-python/9781785286315>

B2. BENEFITS OF PYTHON/JUPYTER

Benefits of Python. I have chosen to use Python and the Jupyter notebook to complete this analysis. Python has all of the plotting (matplotlib and seaborn) and data manipulation packages (numpy, pandas and scipy) that are straight-forward and very easy to use. Jupyter notebooks allows segmented code-execution and the ability to document the analysis using markdown html code. Lastly, the logictics regression analysis required for this analysis is available in Python using the sklearn package.

B3. WHY LOGISTIC REGRESSION

Logistic regression is useful and necessary when the dependent variable is categorical. In our case, the dependent variable is **Churn** which has two (2) unique values, yes and no. Logistic regression can effectively predict the probability of getting the desired outcome based on the set of predictor variables, based on the analysis meeting all of the assumptions stated above.

PART II: EXPLORATORY DATA ANALYSIS

C1. DESCRIBE DATA ANALYSIS PREP AND EXPLORE

Select Data. From the original data, determine which attributes fit the best for the primary research question. Load the data from the provided .csv file as a pandas dataframe.

Mitigate Missing Data. Look through data for missing rows or columns. Also, look for Null or NaN values. If found, decide how best to mitigate the issue.

Remove Data. Once data is determined not to be of value to the analysis, use the pandas .drop() method to remove the data.

Convert Categorical Data. In order to use categorical data in the regression model, each variable must be converted into numeric dummy data. I will use pandas .get_dummies() method. This will generate new numeric variables based on the unique values and this will also remove the original attribute.

Explore Data. Explore customer data by calculating traditional statistics. Look for patterns and relationships between attributes. If possible, create visualizations to add in the exploratory process.

Visualize Data. Continue to explore data and their relationships using histogram, countplots, barplots and scatter plot diagrams. Use matplotlib and sns packages to generate these univariate and bivariate diagrams.

C2-C4. PREPARE AND EXPLORE DATA

Import Packages. Import and configured required math, plotting and model packages.

```
In [2]: # import standard libraries
import numpy as np
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
from IPython.core.display import HTML
from IPython.display import display
```

```
In [3]: # import and configure matplotlib
import matplotlib.pyplot as plt
plt.rc("font", size=14)
```

```
In [4]: # import and configure sklearn
from sklearn.metrics import confusion_matrix
from sklearn import preprocessing
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import classification_report
from sklearn import metrics
```

```
In [5]: # import and configure seaborn
import seaborn as sns
sns.set(style="white")
sns.set(style="whitegrid", color_codes=True)
```

```
In [6]: # import and configure pandas
import pandas as pd
pd.set_option('precision',3)
pd.set_option('max_columns',9)
pd.set_option('display.width', None)
```

Configure Scrollbars. Disable scrollbars in notebook. And, Disable automatically scroll to bottom.

```
In [7]: %%javascript
IPython.OutputArea.prototype._should_scroll = function(lines) {
    return false;
}
```

```
In [8]: %%javascript
require("notebook/js/notebook").Notebook.prototype.scroll_to_bottom = function () {}
```

Stackoverflow (2021, October). Disable iPython Notebook Autoscrolling. Retrieved from: <https://stackoverflow.com/questions/36757301/disable-ipython-notebook-autoscrolling>

Toggle Warnings. Use the following code to toggle warning messages in the notebook. Another piece of code courtesy of stackoverflow (2021).

```
In [9]: from IPython.display import HTML
HTML('''<script>
code_show_err=false;
function code_toggle_err() {
  if (code_show_err){
    $('div.output_stderr').hide();
  } else {
    $('div.output_stderr').show();
  }
  code_show_err = !code_show_err
}
$( document ).ready(code_toggle_err);
</script>
To toggle on/off output_stderr, click <a href="javascript:code_toggle_err()">here</a>.''' )
```

Out[9]: To toggle on/off output_stderr, click [here](#).

Stackoverflow (2021, October). Hide all warnings in ipython. Retrieved from:

<https://stackoverflow.com/questions/9031783/hide-all-warnings-in-ipython>

Helper Functions. Here are some helper functions that will be used throughout the notebook. The coorelation matrix helpers were developed courtesy of stackoverflow (2021).

```

In [10]: def get_redundant_pairs(df):
'''Get diagonal and lower triangular pairs of correlation matrix'''
pairs_to_drop = set()
cols = df.columns
for i in range(0, df.shape[1]):
    for j in range(0, i+1):
        pairs_to_drop.add((cols[i], cols[j]))
return pairs_to_drop

def get_top_abs_correlations(df, n=5):
au_corr = df.corr().abs().unstack()
labels_to_drop = get_redundant_pairs(df)
au_corr = au_corr.drop(labels=labels_to_drop).sort_values(ascending=False)
return au_corr[0:n]

def custom_corr_matrix(df, title):
fig = plt.figure(figsize=(30, 30))
sns.set(font_scale=1.0)
sns.heatmap(data=df.corr().round(1), annot=True, annot_kws={'size':30})
print(get_top_abs_correlations(df))
#plt.savefig('output/' + COURSE + '/fig_corr_matrix_' + title + '.png', facecolor='w')
plt.show()

def plot_histogram(c):
df_yes = df[df.Churn_Yes==1][c]
df_no = df[df.Churn_Yes==0][c]
yes_mean = df_yes.mean();
no_mean = df_no.mean();
fig,ax = plt.subplots(figsize=(6,6))
ax.hist([df_yes,df_no], bins=5, stacked=True)
ax.legend(['Churn - Yes','Churn - No'])
ymin, ymax = ax.get_ylim();
xmin, xmax = ax.get_xlim()
ax.axvline(yes_mean, color='blue', lw=2) # yes mean
ax.axvline(no_mean, color='orangered', lw=2) # no mean
ax.text((xmax-xmin)/2,
        (ymax-ymin)/2,
        'Delta:\n' + str(round(abs(yes_mean - no_mean),2)),
        bbox={'facecolor':'white'})
plt.title('Histogram with target overlay by ' + str(c))
plt.xlabel(c);
plt.ylabel('# Churn');
plt.show();

# helper function to plot grouped bar plot
def plot_stacked(c):
df.groupby([c,target]).size().unstack().plot(kind='bar', stacked=True)

```

Stackoverflow (2021, October). List Highest Correlation Pairs from a Large Correlation Matrix in Pandas. Retrieved from:
<https://stackoverflow.com/questions/17778394/list-highest-correlation-pairs-from-a-large-correlation-matrix-in-pandas>

Constants. Here are a couple of global variables that will be reused throughout the notebook.

```
In [11]: # constants
COURSE = 'd208' # name of course to be added to filename of generated figures and tables.
target = 'Churn' # this is the column name of the primary research column
```

Select Data. The customer dataset as a .csv file is loaded into Python as a Pandas dataframe using the .read_csv() method. After the dataframe is created, I use the df.shape function to show number of rows and columns. To begin the analysis, I have selected to load all of the data from the .csv file.

```
In [12]: # read csv file
import os
df = pd.read_csv(os.path.join('data', 'churn_clean.csv'), header=0)
df.shape
```

Out[12]: (10000, 50)

There are 10,000 customer records with fifty (50) attributes for each customer.

Mitigate Missing Data. Use .info() and .isna().any() methods to view a summary of possible missing data. I do not expect to find any missing data as the dataset provided has already been cleaned.

```
In [13]: # explore missing data
missing = df[df.columns[df.isna().any()]].columns
df_missing = df[missing]
print(df_missing.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Empty DataFrame
```

Analysis of the raw data shows no missing data, each attribute has 10,000 non-null values.

Duplicate Data. Look for duplicate data in rows and columns. This dataset had been provided to this assignment in a very clean, ready state, so I don't expect to find anything here.

```
In [14]: # look for duplicate data - Looking for zero rows
df[df.duplicated()]
```

```
Out[14]:
```

CaseOrder	Customer_id	Interaction	UID	...	Item5	Item6	Item7	Item8
0 rows × 50 columns								

```
In [15]: # check if any cols are duplicated - Looking for False
df.columns.duplicated().any()
```

```
Out[15]: False
```

```
In [16]: # check if any rows are duplicated - Looking for False
df.duplicated().any()
```

```
Out[16]: False
```

Remove Data. Identify columns that are not needed for the analysis and then use the `.drop()` method to remove the data. Looking at the data, I select some of the demographic data, customer identification data and the survey data to be removed.

```
In [17]: # drop unwanted data
cols_to_be_removed = ['City', 'County', 'Zip', 'Job', 'TimeZone', 'State',
                      'Lat', 'Lng', 'UID', 'Customer_id', 'Interaction', 'CaseOrder',
                      'Item1', 'Item2', 'Item3', 'Item4', 'Item5', 'Item6', 'Item7', 'Item8']

# print list of dropped data
print('data to be removed: {}'.format(cols_to_be_removed))

# loop through list, if in current df, drop col
for c in cols_to_be_removed:
    if c in df.columns:
        df.drop(columns = c, inplace=True)
        print('Data named [{}] has been removed.'.format(c))

data to be removed: ['City', 'County', 'Zip', 'Job', 'TimeZone', 'State', 'Lat', 'Lng',
'UID', 'Customer_id', 'Interaction', 'CaseOrder', 'Item1', 'Item2', 'Item3', 'Item4', 'Item5', 'Item6', 'Item7', 'Item8']
Data named [City] has been removed.
Data named [County] has been removed.
Data named [Zip] has been removed.
Data named [Job] has been removed.
Data named [TimeZone] has been removed.
Data named [State] has been removed.
Data named [Lat] has been removed.
Data named [Lng] has been removed.
Data named [UID] has been removed.
Data named [Customer_id] has been removed.
Data named [Interaction] has been removed.
Data named [CaseOrder] has been removed.
Data named [Item1] has been removed.
Data named [Item2] has been removed.
Data named [Item3] has been removed.
Data named [Item4] has been removed.
Data named [Item5] has been removed.
Data named [Item6] has been removed.
Data named [Item7] has been removed.
Data named [Item8] has been removed.
```

Input Variables. Excluding target data, here is the final list of input variables:

Note. Independent variables are sometimes called by different names, they are synonymous, they can be referred to as independent variables, predictor variables, input variables and sometimes, as features.

```
In [18]: # print out and describe input variables
for idx, c in enumerate(df.loc[:, df.columns != target]):
    if df.dtypes[c] == "object":
        print('\n{}. {} is categorical: {}'.format(idx+1, c, df[c].unique()))
        #for idx, name in enumerate(df[c].value_counts().index.tolist()):
        #    print('\t{:<20}:{:>6}'.format(name, df[c].value_counts()[idx]))
        #print('{}'.format(df[c].describe()))
    else:
        print('\n{}. {} is numerical.'.format(idx+1, c))
        #print('{}'.format(df[c].describe().round(3)))
        #groups = df.groupby([target, pd.cut(df[c], bins=4)])
        #print(groups.size().unstack().T)
```

1. Population is numerical.
2. Area is categorical: ['Urban' 'Suburban' 'Rural'].
3. Children is numerical.
4. Age is numerical.
5. Income is numerical.
6. Marital is categorical: ['Widowed' 'Married' 'Separated' 'Never Married' 'Divorce d'].
7. Gender is categorical: ['Male' 'Female' 'Nonbinary'].
8. Outage_sec_perweek is numerical.
9. Email is numerical.
10. Contacts is numerical.
11. Yearly_equip_failure is numerical.
12. Techie is categorical: ['No' 'Yes'].
13. Contract is categorical: ['One year' 'Month-to-month' 'Two Year'].
14. Port_modem is categorical: ['Yes' 'No'].
15. Tablet is categorical: ['Yes' 'No'].
16. InternetService is categorical: ['Fiber Optic' 'DSL' 'None'].
17. Phone is categorical: ['Yes' 'No'].
18. Multiple is categorical: ['No' 'Yes'].
19. OnlineSecurity is categorical: ['Yes' 'No'].
20. OnlineBackup is categorical: ['Yes' 'No'].
21. DeviceProtection is categorical: ['No' 'Yes'].
22. TechSupport is categorical: ['No' 'Yes'].

23. StreamingTV is categorical: ['No' 'Yes'].
24. StreamingMovies is categorical: ['Yes' 'No'].
25. PaperlessBilling is categorical: ['Yes' 'No'].
26. PaymentMethod is categorical: ['Credit Card (automatic)' 'Bank Transfer(automatic)' 'Mailed Check' 'Electronic Check'].
27. Tenure is numerical.
28. MonthlyCharge is numerical.
29. Bandwidth_GB_Year is numerical.

Target Variable. Here is the target variable:

```
In [19]: # print out and describe target variable
for idx, c in enumerate(df.loc[:, df.columns == target]):
    if df.dtypes[c] == "object":
        print('\n{}. {} is categorical: {}'.format(idx+1,c,df[c].unique()))
        for idx,name in enumerate(df[c].value_counts().index.tolist()):
            print('\t{:<8}:{:>6}'.format(name,df[c].value_counts()[idx]))
    else:
        print('\n{}. {} is numerical.'.format(idx+1, c))
```

1. Churn is categorical: ['No' 'Yes'].

No	: 7350
Yes	: 2650

PaymentMethod. The PaymentMethod column has many categories and we can reduce the number of unique values in order to produce a better model. Let's combine all of the data into two (2) categories, 'Automatic' and 'Check'.

```
In [20]: # re-cateogize Marital data
df['PaymentMethod']=np.where(df['PaymentMethod'] == 'Credit Card (automatic)', 'Automatic',
df['PaymentMethod']=np.where(df['PaymentMethod'] == 'Bank Transfer(automatic)', 'Automatic',
df['PaymentMethod']=np.where(df['PaymentMethod'] == 'Mailed Check', 'Check',df['PaymentMethod']
df['PaymentMethod']=np.where(df['PaymentMethod'] == 'Electronic Check', 'Check',df['PaymentMethod'])
```

```
In [21]: # show unique values after grouping
df['PaymentMethod'].unique()
```

```
Out[21]: array(['Automatic', 'Check'], dtype=object)
```

Marital. The Marital column has many categories and we can reduce the number of unique values in order to produce a better model.

Let's combine all of the data into two (2) categories, 'Married' and 'Not_Married'.

```
In [22]: # re-cateogize Marital data
df['Marital']=np.where(df['Marital'] == 'Widowed', 'Not_Married',df['Marital'])
df['Marital']=np.where(df['Marital'] == 'Separated', 'Not_Married',df['Marital'])
df['Marital']=np.where(df['Marital'] == 'Never Married', 'Not_Married',df['Marital'])
df['Marital']=np.where(df['Marital'] == 'Divorced', 'Not_Married',df['Marital'])
```

```
In [23]: # show unique values after grouping
df['Marital'].unique()
```

```
Out[23]: array(['Not_Married', 'Married'], dtype=object)
```

Numeric vs Cateogrical Data. The analysis will use the following variables to separate the numeric and categorical data.

```
In [24]: # variable for numeric data
num_cols = df.select_dtypes(include="number").columns
print(num_cols)

Index(['Population', 'Children', 'Age', 'Income', 'Outage_sec_perweek',
      'Email', 'Contacts', 'Yearly_equip_failure', 'Tenure', 'MonthlyCharge',
      'Bandwidth_GB_Year'],
      dtype='object')
```

```
In [25]: # variable for categorical data
cat_cols = df.select_dtypes(include="object").columns
print(cat_cols)


Index(['Area', 'Marital', 'Gender', 'Churn', 'Techie', 'Contract',
      'Port_modem', 'Tablet', 'InternetService', 'Phone', 'Multiple',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'PaperlessBilling', 'PaymentMethod'],
      dtype='object')
```

CX. EXPLORE TARGET DATA

Explore Target Data. Display a data table and count plot of the target data. Also, calculate percentages of churn and not churned customers in order to determine if the data is balanced or not.

```
In [26]: print(df[target].value_counts())
sns.countplot(x=target, data=df, palette='hls')
plt.show()
```

```
No      7350
Yes      2650
Name: Churn, dtype: int64
```



```
In [27]: # calculate balance
count_no_churn = len(df[df[target]=='No'])
count_churn = len(df[df[target]=='Yes'])
pct_of_no_churn = count_no_churn/(count_no_churn+count_churn)
pct_of_churn = count_churn/(count_no_churn+count_churn)
print('% of customers that did not churn: {:.1%}'.format(pct_of_no_churn ))
print('% of customers that did churn: {:.1%}'.format(pct_of_churn ))
```

```
% of customers that did not churn: 73.5%
% of customers that did churn: 26.5%
```

Unbalanced Data. Observe that the target data is not balanced.

```
In [28]: # describe numerical mean data compared to target
df.groupby(target).mean().round(2).T
```

Out[28]:

	Churn	No	Yes
Population	9830.51	9551.46	
Children	2.09	2.07	
Age	53.01	53.27	
Income	39706.40	40085.76	
Outage_sec_perweek	10.00	10.00	
Email	11.99	12.08	
Contacts	0.99	1.01	
Yearly_equip_failure	0.40	0.38	
Tenure	42.23	13.15	
MonthlyCharge	163.01	199.30	
Bandwidth_GB_Year	3971.86	1785.01	

Observations. Observations:

- A. The **Income** of churned customers is slightly higher.
- B. The **Tenure** of churned customers is significantly lower.
- C. The **MonthlyCharge** of churned customers is considerable higher.
- D. The **Bandwidth_GB_Year** is significantly lower.



















CX. EXPLORE PREDICTOR DATA

Explore Categorical Data. Prior to converting the categorical data for use in the model, as part of exploratory data analysis, I will visualize the original categorical data using a countplot. In a moment, the categorical data will be converted to dummy data and I will lose the original data.


```
In [29]: # plot categorical data - before it gets converted
fig = plt.figure(figsize=(10, 20))

for i, col in enumerate(cat_cols):
    if col != target:
        plt.subplot(10, 3, i+1)
        ax = sns.countplot(y=col, data=df)
        fig.tight_layout(h_pad=4, w_pad=4)

plt.title('Categorical Data')
plt.show()
```



Explore Categorical Data - Means By Unique Value. Show mean values for all

categorical unique values compared to other numeric data.
Explore mean data for each of the categorical variables using
pandas' function `.groupby().mean()`.

In [30]: *# print out mean values of numeric data for a given variable*

```
for c in cat_cols:
    if c != target:
        print('\n\n=====')
        print('\t{}'.format(c.upper()))
        print('=====')
        print(df.groupby(c).mean().round(2).T)
```

```
=====
                AREA
=====
Area              Rural  Suburban  Urban
Population      9732.35  9993.76  9542.22
Children         2.11    2.08    2.07
Age              52.71    53.23    53.30
Income           39667.13 39909.15 39843.92
Outage_sec_perweek 9.99    10.03    9.99
Email            12.05    11.98    12.02
Contacts         1.00    0.99    1.00
Yearly equip_failure 0.40    0.40    0.39
Tenure           35.14    34.37    34.07
MonthlyCharge    172.49    172.49    172.90
Bandwidth_GB_Year 3442.81  3380.28  3354.00
```

```
=====
                MARITAL
=====
Marital          Married  Not_Married
Population      9743.90    9759.55
Children         2.11    2.08
Age              52.84    53.13
Income           40330.36  39683.27
Outage_sec_perweek 10.10    9.98
Email            12.07    12.00
Contacts         1.00    0.99
Yearly equip_failure 0.40    0.40
Tenure           34.81    34.46
MonthlyCharge    172.06    172.76
Bandwidth_GB_Year 3415.47  3386.88
```

```
=====
                GENDER
=====
Gender           Female    Male  Nonbinary
Population      9894.48  9640.72  9135.27
Children         2.08    2.10    2.12
Age              53.09    53.19    50.51
Income           40422.68 39098.45 40962.16
Outage_sec_perweek 9.97    10.04    9.96
Email            11.98    12.06    11.85
Contacts         0.99    1.00    1.00
Yearly equip_failure 0.39    0.40    0.50
Tenure           34.90    34.22    32.65
MonthlyCharge    172.22    173.03    173.13
Bandwidth_GB_Year 3386.78  3407.81  3195.76
```

=====		
TECHIE		
=====		
Techie	No	Yes
Population	9831.00	9387.64
Children	2.09	2.05
Age	53.15	52.73
Income	39738.00	40148.54
Outage_sec_perweek	10.01	9.96
Email	12.04	11.92
Contacts	0.99	1.01
Yearly equip_failure	0.40	0.39
Tenure	34.65	33.94
MonthlyCharge	172.49	173.31
Bandwidth_GB_Year	3401.62	3346.35

=====			
CONTRACT			
=====			
Contract	Month-to-month	One year	Two Year
Population	9664.22	9450.32	10226.47
Children	2.04	2.13	2.15
Age	53.28	52.87	52.81
Income	39848.09	39928.54	39610.28
Outage_sec_perweek	9.98	9.97	10.08
Email	12.04	11.94	12.04
Contacts	0.99	0.99	1.01
Yearly equip_failure	0.40	0.41	0.39
Tenure	34.24	34.21	35.43
MonthlyCharge	172.14	173.67	172.80
Bandwidth_GB_Year	3363.43	3375.53	3471.40

=====		
PORT_MODEM		
=====		
Port_modem	No	Yes
Population	9636.82	9884.53
Children	2.06	2.11
Age	52.96	53.21
Income	40215.46	39370.33
Outage_sec_perweek	9.98	10.02
Email	11.97	12.06
Contacts	1.00	0.99
Yearly equip_failure	0.39	0.40
Tenure	34.30	34.77
MonthlyCharge	172.62	172.63
Bandwidth_GB_Year	3373.98	3411.97

=====		
TABLET		
=====		
Tablet	No	Yes
Population	9745.02	9783.62
Children	2.09	2.09

Age	53.14	52.94
Income	39696.41	40065.90
Outage_sec_perweek	9.99	10.03
Email	12.03	11.98
Contacts	1.00	0.98
Yearly_equip_failure	0.40	0.40
Tenure	34.52	34.55
MonthlyCharge	172.42	173.10
Bandwidth_GB_Year	3389.30	3399.46

=====

INTERNETSERVICE

=====

InternetService	DSL	Fiber Optic	None
Population	9860.59	9751.87	9597.06
Children	2.08	2.08	2.12
Age	52.95	53.11	53.23
Income	40535.84	39267.31	39738.54
Outage_sec_perweek	9.96	10.02	10.05
Email	12.01	12.03	12.01
Contacts	1.01	1.00	0.96
Yearly_equip_failure	0.40	0.39	0.40
Tenure	34.99	34.41	34.01
MonthlyCharge	167.06	186.49	152.98
Bandwidth_GB_Year	3701.97	3239.85	3204.42

=====

PHONE

=====

Phone	No	Yes
Population	9387.80	9794.51
Children	2.10	2.09
Age	52.46	53.14
Income	39954.32	39791.76
Outage_sec_perweek	10.10	9.99
Email	12.08	12.01
Contacts	0.98	1.00
Yearly_equip_failure	0.41	0.40
Tenure	34.26	34.55
MonthlyCharge	175.17	172.36
Bandwidth_GB_Year	3383.27	3393.27

=====

MULTIPLE

=====

Multiple	No	Yes
Population	9773.12	9737.19
Children	2.13	2.04
Age	53.17	52.97
Income	39812.52	39800.39
Outage_sec_perweek	9.97	10.04
Email	12.02	12.01
Contacts	1.01	0.98
Yearly_equip_failure	0.40	0.40
Tenure	34.78	34.23
MonthlyCharge	157.30	190.55

Bandwidth_GB_Year	3378.56	3408.47
-------------------	---------	---------

=====

ONLINESECURITY

=====

OnlineSecurity	No	Yes
Population	9621.44	9999.31
Children	2.08	2.11
Age	53.25	52.77
Income	40025.41	39414.45
Outage_sec_perweek	10.01	9.99
Email	12.07	11.91
Contacts	0.99	1.01
Yearly equip_failure	0.41	0.38
Tenure	34.48	34.62
MonthlyCharge	171.10	175.37
Bandwidth_GB_Year	3358.09	3453.86

=====

ONLINEBACKUP

=====

OnlineBackup	No	Yes
Population	9621.26	9921.53
Children	2.09	2.09
Age	53.02	53.15
Income	39959.49	39620.91
Outage_sec_perweek	10.00	10.00
Email	12.05	11.98
Contacts	0.98	1.01
Yearly equip_failure	0.40	0.39
Tenure	34.03	35.13
MonthlyCharge	162.54	184.93
Bandwidth_GB_Year	3309.74	3493.05

=====

DEVICEPROTECTION

=====

DeviceProtection	No	Yes
Population	9804.98	9694.59
Children	2.08	2.10
Age	52.81	53.42
Income	39456.16	40255.90
Outage_sec_perweek	9.94	10.08
Email	12.01	12.02
Contacts	0.99	0.99
Yearly equip_failure	0.40	0.39
Tenure	35.18	33.69
MonthlyCharge	166.45	180.53
Bandwidth_GB_Year	3407.52	3372.92

=====

TECHSUPPORT

=====

TechSupport	No	Yes
Population	9830.41	9633.48

Children	2.09	2.09
Age	52.80	53.54
Income	39613.95	40128.55
Outage_sec_perweek	10.03	9.95
Email	11.96	12.10
Contacts	1.00	0.98
Yearly_equip_failure	0.40	0.40
Tenure	34.55	34.48
MonthlyCharge	168.62	179.29
Bandwidth_GB_Year	3392.14	3392.68

=====

STREAMINGTV

=====

StreamingTV	No	Yes
Population	9850.32	9660.10
Children	2.10	2.07
Age	53.08	53.07
Income	39899.33	39711.86
Outage_sec_perweek	9.98	10.03
Email	12.00	12.03
Contacts	0.99	0.99
Yearly_equip_failure	0.40	0.40
Tenure	34.46	34.59
MonthlyCharge	152.21	193.63
Bandwidth_GB_Year	3275.33	3512.73

=====

STREAMINGMOVIES

=====

StreamingMovies	No	Yes
Population	9839.61	9669.78
Children	2.08	2.10
Age	52.87	53.29
Income	39867.22	39743.92
Outage_sec_perweek	9.95	10.06
Email	12.01	12.02
Contacts	0.98	1.01
Yearly_equip_failure	0.40	0.40
Tenure	34.59	34.46
MonthlyCharge	147.08	199.32
Bandwidth_GB_Year	3294.87	3494.20

=====

PAPERLESSBILLING

=====

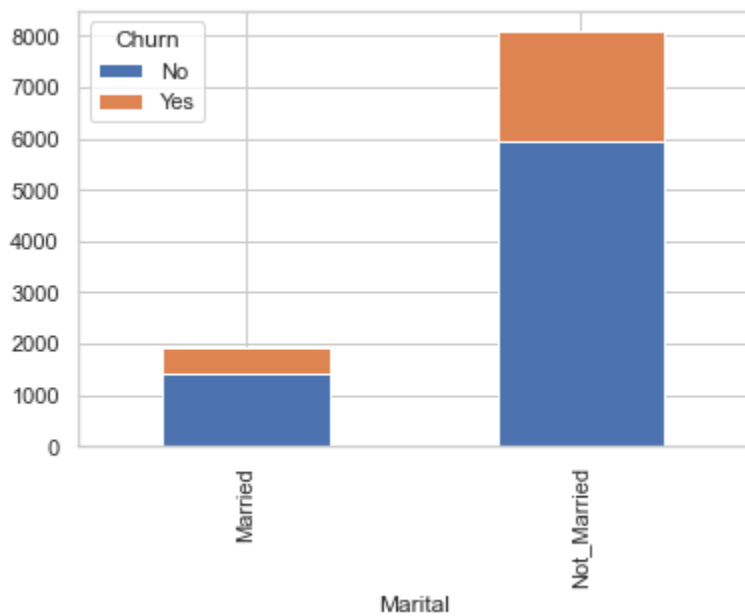

PaperlessBilling	No	Yes
Population	9607.27	9861.09
Children	2.07	2.10
Age	53.02	53.12
Income	40212.98	39522.65
Outage_sec_perweek	10.04	9.97
Email	12.06	11.98
Contacts	1.00	0.99
Yearly_equip_failure	0.39	0.41
Tenure	34.55	34.51

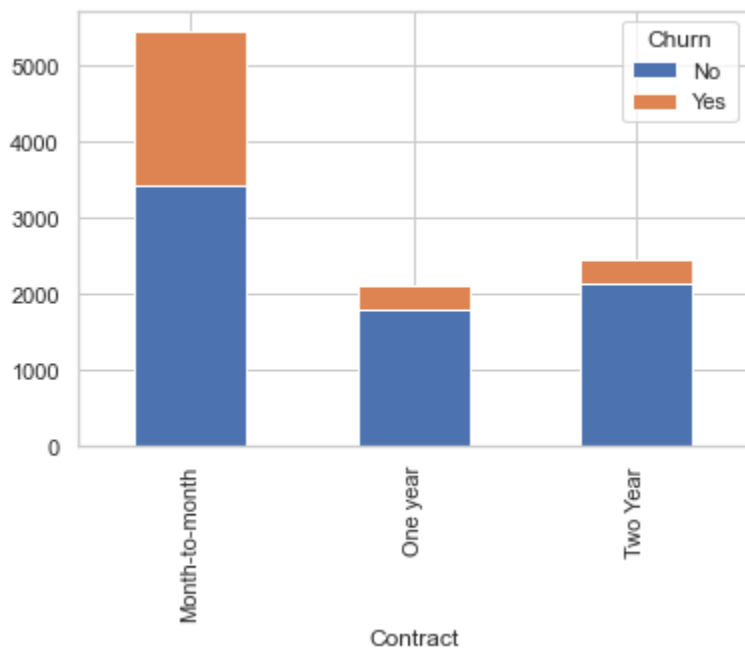
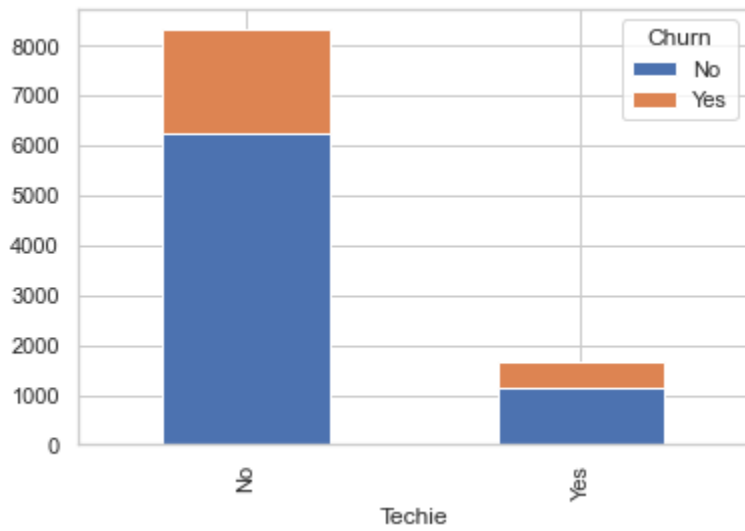

MonthlyCharge	172.56	172.67
Bandwidth_GB_Year	3398.79	3387.82




```
=====
PAYMENTMETHOD
=====
PaymentMethod      Automatic      Check
Population          9887.14      9657.58
Children            2.09         2.09
Age                 53.48        52.77
Income              39296.10     40194.18
Outage_sec_perweek  10.01         9.99
Email               12.04         12.00
Contacts            1.02          0.98
Yearly_equip_failure 0.40          0.40
Tenure              34.45         34.58
MonthlyCharge       173.24        172.16
Bandwidth_GB_Year   3385.03       3397.88
```


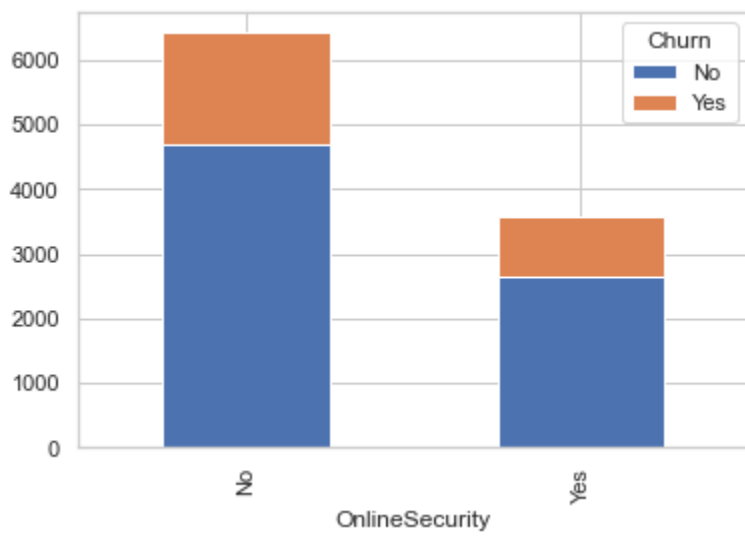
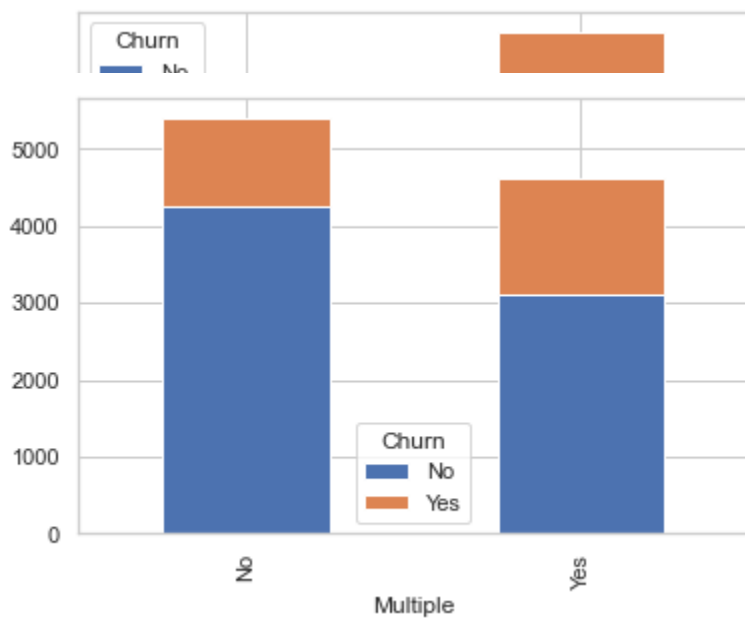
Explore Categorical Data - Grouped Bar Plot with Target Overlay. Grouped bar plot of each categorical data with target data overlaid.

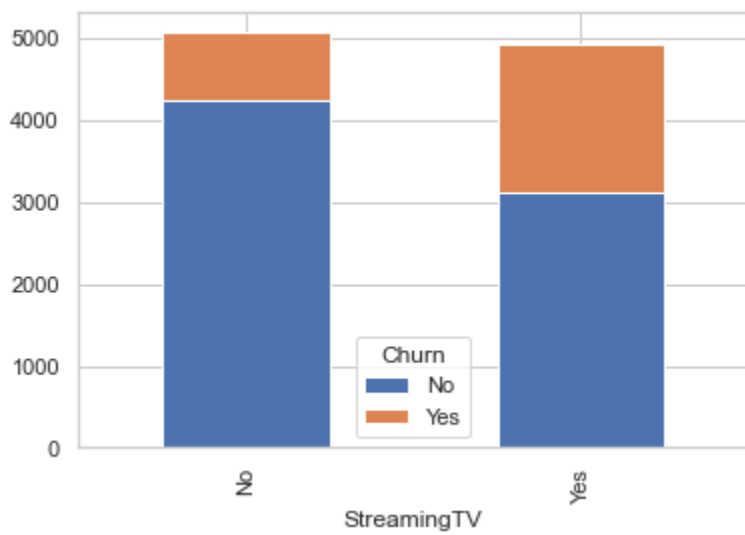
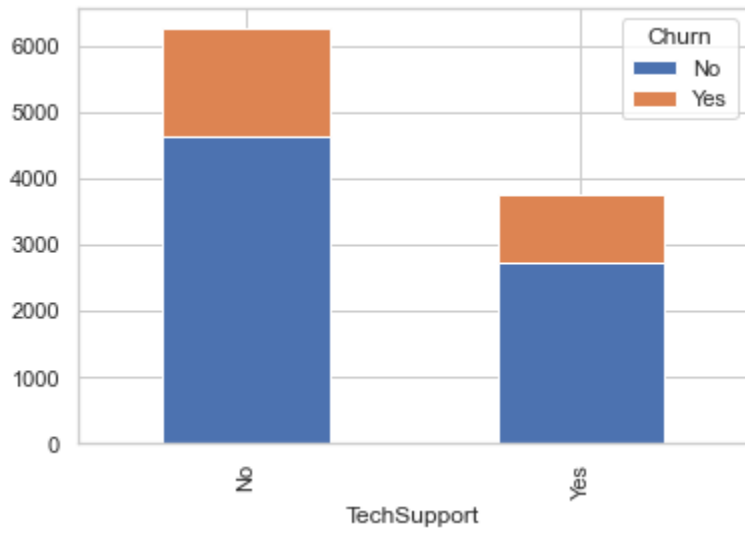

```
In [31]: # plot each variable vs. target overlay
for c in cat_cols:
    if c != target:
        plot_stacked(c)
```
















Good Predictors. Based on the plots, it appears that the following variables might be good predictors of the outcome:

- A. Marital
- B. Gender
- C. Techie
- D. Contract
- E. Tablet
- F. InternetService
- G. OnlineSecurity - maybe
- H. OnlineBackup - maybe
- I. TechSupport - maybe
- J. StreamingTV
- K. StreamingMovies
- L. PaperlessBilling
- M. PaymentMethod

Convert Selected Categorical Data. Now that I have selected some of the categorical data that seem to be a good predictors of the outcome, I will convert these categorical data to dummy, numeric data. Each new variable will have a value of either one (1) or zero (0).

```
In [32]: # convert categorical data
for c in cat_cols:
    if c in df.columns:
        df = pd.get_dummies(df, columns=[c], drop_first=True)
pred_vars = df.select_dtypes(include="uint8").columns.tolist()
print(pred_vars)
```

```
['Area_Suburban', 'Area_Urban', 'Marital_Not_Married', 'Gender_Male', 'Gender_Nonbinary',
'Churn_Yes', 'Techie_Yes', 'Contract_One year', 'Contract_Two Year', 'Port_modem_Yes', 'T
ablet_Yes', 'InternetService_Fiber Optic', 'InternetService_None', 'Phone_Yes', 'Multiple
_Yes', 'OnlineSecurity_Yes', 'OnlineBackup_Yes', 'DeviceProtection_Yes', 'TechSupport_Ye
s', 'StreamingTV_Yes', 'StreamingMovies_Yes', 'PaperlessBilling_Yes', 'PaymentMethod_Chec
k']
```

```
In [33]: # reset the global target variable using its dummy variable
target = 'Churn_Yes'
```

Describe Numeric Data. Traditional statistics for numeric data.


```
In [34]: # describe numeric data
df[num_cols].describe().round(3).T
```

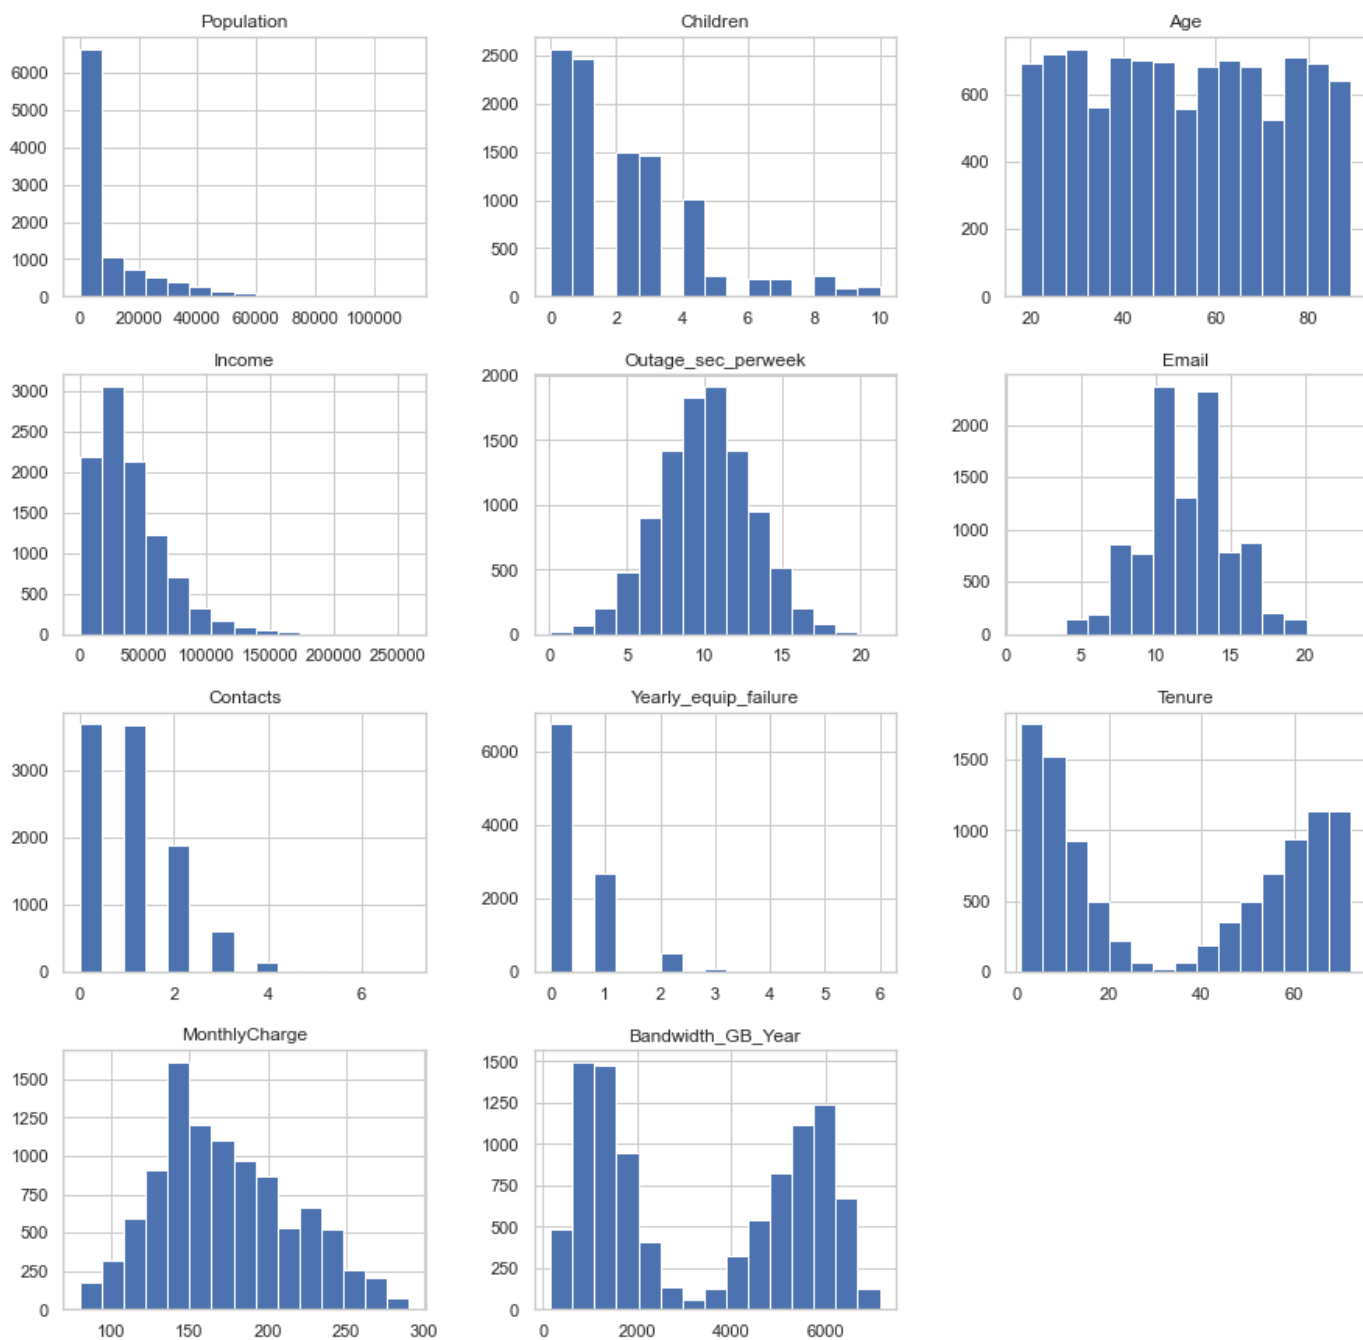
Out[34]:

	count	mean	std	min	25%	50%	75%	max
Population	10000.0	9756.562	14432.699	0.000	738.000	2910.500	13168.000	111850.000
Children	10000.0	2.088	2.147	0.000	0.000	1.000	3.000	10.000
Age	10000.0	53.078	20.699	18.000	35.000	53.000	71.000	89.000
Income	10000.0	39806.927	28199.917	348.670	19224.718	33170.605	53246.170	258900.700
Outage_sec_perweek	10000.0	10.002	2.976	0.100	8.018	10.019	11.969	21.207
Email	10000.0	12.016	3.026	1.000	10.000	12.000	14.000	23.000
Contacts	10000.0	0.994	0.988	0.000	0.000	1.000	2.000	7.000
Yearly_equip_failure	10000.0	0.398	0.636	0.000	0.000	0.000	1.000	6.000
Tenure	10000.0	34.526	26.443	1.000	7.918	35.431	61.480	71.999
MonthlyCharge	10000.0	172.625	42.943	79.979	139.979	167.485	200.735	290.160
Bandwidth_GB_Year	10000.0	3392.342	2185.295	155.507	1236.471	3279.537	5586.141	7158.982

Univariate Histogram Plot of Numeric Predictor Data. Here are the histogram plots for numeric data.


```
In [35]: # histogram plot numeric data
fig = plt.figure(figsize=(10, 20))
ax = df[num_cols].hist(bins = 15, figsize=(15,15))
plt.title('Numeric Data')
fig.tight_layout(h_pad=5, w_pad=5)
plt.show()
```

<Figure size 720x1440 with 0 Axes>



Bivariate Barplot of Numeric Predictor Data. Here are the histogram plots for the numeric data. Each of these plots show the frequency counts based on the yes or no value of the target variable.


```
In [36]: # create histogram with target overlay  
plot_histogram('MonthlyCharge')
```



Observations. Observations:

A. The mean value for churned customers is 36.29 units higher.

```
In [37]: # create histogram with target overlay
plot_histogram('Tenure')
```




Observations.

Observations:

- A. The mean value for churned customers is 29.09 units lower.

```
In [38]: # create histogram with target overlay
plot_histogram('Bandwidth_GB_Year')
```




Observations.

Observations:

A. The mean value for churned customers is 2186.85 units lower.

```
In [39]: # create histogram with target overlay  
plot_histogram('Age')
```




Observations.

Observations:

- A. There appears to be no distinction between churned and non-churned customers.


```
In [40]: # create histogram with target overlay  
plot_histogram('Income')
```



Observations. Observations :

A. There appears to be no distinction between churned and non-churned customers.


```
In [41]: # create histogram with target overlay  
plot_histogram('Outage_sec_perweek')
```



Observations. Observations:

A. There appears to be no distinction between churned and non-churned customers.



```
In [42]: # create histogram with target overlay  
plot_histogram('Email')
```



Observations. Observations:

A. There appears to be no distinction between churned and non-churned customers.

```
In [43]: # create histogram with target overlay  
plot_histogram('Contacts')
```




Observations.

Observations :

- A. There appears to be no distinction between churned and non-churned customers.


```
In [44]: # create histogram with target overlay  
plot_histogram('Yearly equip_failure')
```



Observations. Observations:

A. There appears to be no distinction between churned and non-churned customers.


```
In [45]: # create histogram with target overlay  
plot_histogram('Population')
```



Observations. Observations :

A. There appears to be no distinction between churned and non-churned customers.

```
In [46]: # create histogram with target overlay
plot_histogram('Children')
```



Observations. Observations:

A. There appears to be no distinction between churned and non-churned customers.

Drop Data. Need to get rid of the variable **Churn_No** in order to prevent collinearity with **Churn_Yes**

```
In [47]: # drop unwanted data
cols_to_be_removed = ['Churn_No']

# print list of dropped data
print('data to be removed: {}'.format(cols_to_be_removed))

# Loop through list, if in current df, drop col
for c in cols_to_be_removed:
    if c in df.columns:
        df.drop(columns = c, inplace=True)
        print('Data named [{}] has been removed.'.format(c))
```

data to be removed: ['Churn_No']

C5. PROVIDE COPY OF DATA

Final Data. Here is the final list of columns after all data cleaning.

```
In [48]: # Provide copy of the prepared data set.
final_data = 'd208_task2_final_data.csv'
df.to_csv(final_data, index=False, header=True)
print('File saved to: {}'.format(final_data))
print(df.columns.to_series().groupby(df.dtypes).groups)
```

```
File saved to: d208_task2_final_data.csv
{uint8: ['Area_Suburban', 'Area_Urban', 'Marital_Not_Married', 'Gender_Male', 'Gender_Non
binary', 'Churn_Yes', 'Techie_Yes', 'Contract_One_year', 'Contract_Two_Year', 'Port_modem
_Yes', 'Tablet_Yes', 'InternetService_Fiber_Optic', 'InternetService_None', 'Phone_Yes',
'Multiple_Yes', 'OnlineSecurity_Yes', 'OnlineBackup_Yes', 'DeviceProtection_Yes', 'TechSu
pport_Yes', 'StreamingTV_Yes', 'StreamingMovies_Yes', 'PaperlessBilling_Yes', 'PaymentMet
hod_Check'], int64: ['Population', 'Children', 'Age', 'Email', 'Contacts', 'Yearly equip_
failure'], float64: ['Income', 'Outage_sec_perweek', 'Tenure', 'MonthlyCharge', 'Bandwidt
h_GB_Year']}
```

PART III: MODEL COMPARISON AND ANALYSIS

D1. INITIAL MODEL

Balance Data. We saw earlier that the data is not balanced, so I am using the SMOTE package to oversample the data in order to balance out the data before running the model.

```
In [49]: # rebalance data using SMOTE oversample
X = df.loc[:, df.columns != 'Churn_Yes']
y = df.loc[:, df.columns == 'Churn_Yes']
from imblearn.over_sampling import SMOTE
os = SMOTE(random_state=0)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
columns = X_train.columns
os_data_X, os_data_y = os.fit_resample(X_train, y_train)
os_data_X = pd.DataFrame(data=os_data_X, columns=columns)
os_data_y = pd.DataFrame(data=os_data_y, columns=['Churn_Yes'])
# we can check the numbers of our data
print("length of oversampled data is ", len(os_data_X))
print("Number of no churn in oversampled data", len(os_data_y[os_data_y['Churn_Yes']==0]))
print("Number of churn", len(os_data_y[os_data_y['Churn_Yes']==1]))
print("Proportion of no churn data in oversampled data is ", len(os_data_y[os_data_y['Churn_Yes']==0])/len(os_data_y))
print("Proportion of churn data in oversampled data is ", len(os_data_y[os_data_y['Churn_Yes']==1])/len(os_data_y))
```

```
length of oversampled data is 10298
Number of no churn in oversampled data 5149
Number of churn 5149
Proportion of no churn data in oversampled data is 0.5
Proportion of churn data in oversampled data is 0.5
```

RFE Feature Reduction. RFE in the sklearn package finds and ranks the features with the most potential to the model.

```
In [50]: # RFE feature reduction
data_final_vars=df.columns.values.tolist()
y=[target]
X=[i for i in data_final_vars if i not in y]
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
rfe = RFE(logreg, 20)
rfe = rfe.fit(os_data_X, os_data_y.values.ravel())
features = []
print('The following features are selected:')
for i in range(os_data_X.shape[1]):
    if rfe.support_[i] == True:
        features.append(os_data_X.columns[i])
        print('Column: %d, Rank: %.3f, Feature: %s' %
              (i, rfe.ranking_[i],
               os_data_X.columns[i]))
```

The following features are selected:

```
Column: 4, Rank: 1.000, Feature: Outage_sec_perweek
Column: 5, Rank: 1.000, Feature: Email
Column: 8, Rank: 1.000, Feature: Tenure
Column: 9, Rank: 1.000, Feature: MonthlyCharge
Column: 11, Rank: 1.000, Feature: Area_Suburban
Column: 12, Rank: 1.000, Feature: Area_Urban
Column: 13, Rank: 1.000, Feature: Marital_Not_Married
Column: 15, Rank: 1.000, Feature: Gender_Nonbinary
Column: 16, Rank: 1.000, Feature: Techie_Yes
Column: 17, Rank: 1.000, Feature: Contract_One year
Column: 18, Rank: 1.000, Feature: Contract_Two Year
Column: 20, Rank: 1.000, Feature: Tablet_Yes
Column: 21, Rank: 1.000, Feature: InternetService_Fiber Optic
Column: 22, Rank: 1.000, Feature: InternetService_None
Column: 23, Rank: 1.000, Feature: Phone_Yes
Column: 24, Rank: 1.000, Feature: Multiple_Yes
Column: 25, Rank: 1.000, Feature: OnlineSecurity_Yes
Column: 26, Rank: 1.000, Feature: OnlineBackup_Yes
Column: 27, Rank: 1.000, Feature: DeviceProtection_Yes
Column: 28, Rank: 1.000, Feature: TechSupport_Yes
```

Initial Model. Use the RFE features analysis above to select the best features for the initial iteration of the model.

```
In [51]: # initial model
X=os_data_X[features] # from RFE above
Xc = sm.add_constant(X) # reset
y=os_data_y[target]
logit_model=sm.Logit(y,Xc)
result=logit_model.fit()
print(result.summary2())
```

Optimization terminated successfully.
Current function value: inf
Iterations 9

Results: Logit


```
=====
Model:                               Logit                               Pseudo R-squared:         inf
Dependent Variable:                   Churn_Yes                           AIC:                      inf
Date:                                2021-10-11 03:50                       BIC:                      inf
No. Observations:                     10298                          Log-Likelihood:          -inf
Df Model:                             20                             LL-Null:                  0.0000
Df Residuals:                         10277                          LLR p-value:              1.0000
Converged:                            1.0000                          Scale:                    1.0000
No. Iterations:                       9.0000

-----
              Coef.  Std.Err.    z      P>|z|    [0.025  0.975]
-----
const                -5.1237    0.3713  -13.7980  0.0000   -5.8515  -4.3959
Outage_sec_perweek   -0.0043    0.0150   -0.2860  0.7749   -0.0337  0.0251
Email                -0.0257    0.0147   -1.7475  0.0805   -0.0545  0.0031
Tenure               -0.1302    0.0035  -37.1082  0.0000   -0.1370 -0.1233
MonthlyCharge         0.0738    0.0020   36.3557  0.0000    0.0698  0.0778
Area_Suburban        -0.5842    0.1057   -5.5262  0.0000   -0.7915 -0.3770
Area_Urban           -0.3605    0.1055   -3.4183  0.0006   -0.5672 -0.1538
Marital_Not_Married  -0.2252    0.1042   -2.1617  0.0306   -0.4294 -0.0210
Gender_Nonbinary     -0.2690    0.3109   -0.8652  0.3869   -0.8784  0.3404
Techie_Yes           0.9889    0.1237    7.9952  0.0000    0.7465  1.2313
Contract_One year    -4.0193    0.1574  -25.5387  0.0000   -4.3278 -3.7109
Contract_Two Year    -4.0660    0.1493  -27.2303  0.0000   -4.3587 -3.7734
Tablet_Yes           -0.2900    0.0987   -2.9375  0.0033   -0.4835 -0.0965
InternetService_Fiber Optic -3.1580    0.1188  -26.5901  0.0000   -3.3908 -2.9253
InternetService_None -1.0368    0.1242   -8.3508  0.0000   -1.2802 -0.7935
Phone_Yes            -0.5538    0.1476   -3.7514  0.0002   -0.8431 -0.2644
Multiple_Yes         -0.7568    0.0959   -7.8942  0.0000   -0.9447 -0.5689
OnlineSecurity_Yes   -0.5733    0.0939   -6.1042  0.0000   -0.7573 -0.3892
OnlineBackup_Yes     -0.9050    0.0928   -9.7509  0.0000   -1.0869 -0.7231
DeviceProtection_Yes -0.6084    0.0904   -6.7330  0.0000   -0.7855 -0.4313
TechSupport_Yes      -0.9436    0.0932  -10.1236  0.0000   -1.1262 -0.7609
=====
```

Confusion Matrix - Initial Model. Display confusion matrix of the test and prediction data from the intial model.

In [52]: *# confusion matrix for initial model*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
lgr = LogisticRegression()
lgr.fit(X_train, y_train)
predicted = lgr.predict(X_test)
expected = y_test
confusion = pd.DataFrame(confusion_matrix(y_true=expected, y_pred=predicted),
                          index=range(2), columns=range(2))
axes = sns.heatmap(confusion, annot=True, cmap='nipy_spectral_r', fmt='g')
```



In [53]: *# calculate number and percent of predictions*

```
correct = sum(np.diagonal(confusion)) # on diag
total = confusion.values.sum()
incorrect = total - correct # off diag
print('Correct predictions on diagonal: {} ({:.0%})'.format( correct, correct / total ))
print('Incorrect predictions off diagonal: {} ({:.0%})'.format( incorrect, incorrect / total ))
```

```
Correct predictions on diagonal: 2865 (93%)
Incorrect predictions off diagonal: 225 (7%)
```

High Coorelation. Use coorelation matrix to find predictor pairs with high coorelation.

```
In [54]: # find predictor pairs with high coorelation
#custom_corr_matrix(X, 'Model_2')
get_top_abs_correlations(X, 20)
```

```
Out[54]: Area_Suburban      Area_Urban      0.359
InternetService_Fiber Optic  InternetService_None  0.318
MonthlyCharge      Multiple_Yes      0.287
                  InternetService_None  0.266
Contract_One year   Contract_Two Year   0.187
MonthlyCharge      OnlineBackup_Yes    0.168
                  InternetService_Fiber Optic  0.161
Tenure             Contract_Two Year   0.140
                  OnlineSecurity_Yes    0.108
                  OnlineBackup_Yes    0.106
                  InternetService_Fiber Optic  0.105
                  Contract_One year    0.099
                  Area_Suburban      0.097
Contract_Two Year    InternetService_Fiber Optic  0.096
Tenure             Area_Urban      0.089
Contract_Two Year    OnlineSecurity_Yes    0.086
Tenure             Tablet_Yes      0.086
                  MonthlyCharge      0.082
Contract_One year   Tablet_Yes      0.078
Tenure             InternetService_None  0.077
dtype: float64
```

D3. REDUCED MODEL

Final Model. Remove features with high P-values or high multi-collinearity. It looks like there are a few of the input variables that have high multi-collinearity. Then re-run the model.

```
In [55]: # update model
features.remove('Gender_Nonbinary') # high p-value
features.remove('Outage_sec_perweek') # high p-value
features.remove('Email') # high p-value
features.remove('Marital_Not_Married') # high p-value
features.remove('Area_Urban') # high collinearity
features.remove('InternetService_None') # high collinearity
X=os_data_X[features]
y=os_data_y[target]
Xc = sm.add_constant(X) # reset
logit_model=sm.Logit(y,Xc)
result=logit_model.fit()
print(result.summary2())
```

Optimization terminated successfully.

Current function value: inf

Iterations 9


Results: Logit

```
=====
Model:                Logit                Pseudo R-squared:      inf
Dependent Variable:    Churn_Yes             AIC:                  inf
Date:                  2021-10-11 03:50       BIC:                  inf
No. Observations:      10298                 Log-Likelihood:       -inf
Df Model:              14                    LL-Null:              0.0000
Df Residuals:          10283                 LLR p-value:          1.0000
Converged:              1.0000                Scale:                1.0000
No. Iterations:        9.0000

-----
              Coef.  Std.Err.   z      P>|z|    [0.025  0.975]
-----
const                -6.2088    0.2727 -22.7693 0.0000 -6.7433 -5.6744
Tenure                -0.1294    0.0034 -37.6573 0.0000 -0.1362 -0.1227
MonthlyCharge         0.0753    0.0020  37.5204 0.0000  0.0714  0.0792
Area_Suburban        -0.4517    0.0942 -4.7945 0.0000 -0.6364 -0.2671
Techie_Yes           0.9742    0.1243  7.8385 0.0000  0.7306  1.2178
Contract_One year    -4.0021    0.1559 -25.6762 0.0000 -4.3076 -3.6966
Contract_Two Year    -4.0364    0.1488 -27.1184 0.0000 -4.3281 -3.7447
Tablet_Yes           -0.3568    0.0973 -3.6684 0.0002 -0.5475 -0.1662
InternetService_Fiber Optic -2.8809    0.1088 -26.4682 0.0000 -3.0943 -2.6676
Phone_Yes            -0.5865    0.1434 -4.0909 0.0000 -0.8675 -0.3055
Multiple_Yes         -0.8422    0.0942 -8.9423 0.0000 -1.0268 -0.6576
OnlineSecurity_Yes   -0.6116    0.0929 -6.5859 0.0000 -0.7937 -0.4296
OnlineBackup_Yes     -0.9683    0.0916 -10.5734 0.0000 -1.1478 -0.7888
DeviceProtection_Yes -0.6737    0.0892 -7.5531 0.0000 -0.8485 -0.4989
TechSupport_Yes      -1.0144    0.0921 -11.0129 0.0000 -1.1949 -0.8339
=====
```

Confusion Matrix - Final Model. Display confusion matrix of the test and prediction data.

```
In [56]: # confusion matrix for final model
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
lgr = LogisticRegression()
lgr.fit(X_train, y_train)
predicted = lgr.predict(X_test)
expected = y_test
confusion = pd.DataFrame(confusion_matrix(y_true=expected, y_pred=predicted),
                          index=range(2), columns=range(2))
axes = sns.heatmap(confusion, annot=True, cmap='nipy_spectral_r', fmt='g')
```



```
In [57]: # calculate number and percent of predictions
correct = sum(np.diagonal(confusion)) # on diag
total = confusion.values.sum()
incorrect = total - correct # off diag
print('Correct predictions on diagonal: {} ({:.0%})'.format( correct, correct / total ))
print('Incorrect predictions off diagonal: {} ({:.0%})'.format( incorrect, incorrect / total ))
```

Correct predictions on diagonal: 2863 (93%)
 Incorrect predictions off diagonal: 227 (7%)


Classification Report. Classification report below:

```
In [58]: # classification report
print(classification_report(expected, predicted))
```

	precision	recall	f1-score	support
0	0.93	0.92	0.93	1545
1	0.92	0.94	0.93	1545
accuracy			0.93	3090
macro avg	0.93	0.93	0.93	3090
weighted avg	0.93	0.93	0.93	3090

ROC Curve. ROC curve below:

```
In [59]: # plot ROC Curve
logit_roc_auc = roc_auc_score(y_test, lgr.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, lgr.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' %logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



According to Li (2017), "The receiver operating characteristic (ROC) curve is another common tool used with binary classifiers. The dotted line represents the ROC curve of a purely random classifier; a good classifier stays as far away from that line as possible (toward the topleft corner)." It looks like we have generated a good model.

Li, S. (2017, Sep 28). Building A Logistic Regression in Python, Step by Step.

Retrieved from: <https://towardsdatascience.com/building-a-logistic-regression-in-python-step-by-step-becd4d56c9c8>

E1. EXPLAIN

The features or input variables used for the initial model were selected using the RFE selection process.

Both models have about the same number of correct predictions, about 93%. However, the final model removed some of the input variables that had high multicollinearity and another couple removed because of high p-value. The reduced model is cleaner but has about the same results.

E2. PROVIDE OUTPUT AND CALCULATIONS

Output. All output and calculations are provided within this Jupyter notebook.

E3. PROVIDE CODE

Code. All code is provided within this Jupyter notebook. There is a .PDF file of the final notebook and the .IPYNB file submitted with the assignment which has all of the code input and outputs.

PART IV: DATA SUMMARY AND IMPLICATIONS

F1. DISCUSS RESULTS

F1. Summarize your findings and assumptions by doing the following: 1. Discuss the results of your data analysis, including the following elements: • a regression equation for

the reduced model • an interpretation of coefficients of the statistically significant variables of the model • the statistical and practical significance of the model • the limitations of the data analysis


Regression Equation. Here is the regression equation of the reduced model:

```
In [60]: # equation of the regression line/plane
print('Logit: {:.2f}'.format(logit_roc_auc))
equation = result.summary2().tables[1]
print('Estimate [{}] as L = '.format(result.summary2().tables[0][1][1]))
for i in equation.itertuples():
    print('    {:.3f} x ( {} ) '.format(i[1],i[0]))
```

```
Logit: 0.93
Estimate [Churn_Yes] as L =
    -6.209 x ( const )
    -0.129 x ( Tenure )
    +0.075 x ( MonthlyCharge )
    -0.452 x ( Area_Suburban )
    +0.974 x ( Techie_Yes )
    -4.002 x ( Contract_One year )
    -4.036 x ( Contract_Two Year )
    -0.357 x ( Tablet_Yes )
    -2.881 x ( InternetService_Fiber Optic )
    -0.587 x ( Phone_Yes )
    -0.842 x ( Multiple_Yes )
    -0.612 x ( OnlineSecurity_Yes )
    -0.968 x ( OnlineBackup_Yes )
    -0.674 x ( DeviceProtection_Yes )
    -1.014 x ( TechSupport_Yes )
```

Interpret Logistic Regression Coefficients. Each regression coefficient describes the estimated change in the log-odds of the response variable when the coefficient's predictor variable increases by one. (LaRose, 2019)

```
In [61]: # visualization of the log coef using exp(x) plot
import matplotlib.pyplot as plt
X_coeff = []
for i in equation.itertuples():
    X_coeff.append(i[1])
X_coeff.sort()
x = X_coeff
y = np.exp(x)
plt.figure()
plt.plot(X_coeff, y, color="red", marker="+")
plt.xlim([-7.0, 2.0])
plt.ylim([-1.0, 3.0])
plt.xlabel('$x$')
plt.ylabel('$\exp(x)$')
plt.show()
```



Here are a few specific examples:

Techie_Yes. Consider the binary predictor variable **Techie_Yes**. The regression coefficient is +0.974. By calculating $e^{0.974} = 2.648$, we find that a customer is about 2.6 times as likely to churn if they are a "techie" compared to if they are not.

MonthlyCharge. Consider the binary predictor variable **MonthlyCharge**. The regression coefficient is +0.075. By calculating $e^{0.075} = 1.079$, we find that a customer is about 1.1 times as likely to churn for every 1 unit of increase in **MonthlyCharge**.

Plus/Minus. You see some of the coefficients are positive and some negative, the positive factors indicate higher likelihood of churn and the negative factor indicate lower likelihood of churn. You can see that +0.974 **Techie_Yes** will cancel out with the -0.968 if the customer has the **OnlineBackup_Yes** service.

F2. RECOMMENDATIONS

F2. Recommend a course of action based on your results.

Recommendations. Customers will be less likely to churn if their **MonthlyCharge** is minimized and if the customer has any of the additionally available services. Focus marketing efforts on which additional services are best for each customer, maybe bundle some of the services at a slightly reduced monthly payment. Increase customers' awareness of the value of the additional services for what they are paying each month. Keeping their monthly payment low and increasing the number of extra services will minimize likelihood of churn, and provide company with increase in the customer lifetime revenue.

PART V: DEMONSTRATION

G. VIDEO

G. Provide a Panopto video recording that includes all of the following elements: • a demonstration of the functionality of the code used for the analysis • an identification of the version of the programming environment • a comparison of the two logistic regression models you used in your analysis • an interpretation of the coefficients

Video was created and posted to the WGU class dropbox.

Here is the link

<https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=a90e1811-2440-4e27-a081-adbb01855443>
[_ \(https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=a90e1811-2440-4e27-a081-adbb01855443\)_](https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=a90e1811-2440-4e27-a081-adbb01855443)

In []: